

StaQC: A Systematically Mined Question-Code Dataset from Stack Overflow

Ziyu Yao

The Ohio State University

The Web Conference 2018, May 25th

In collaboration with Prof. Daniel S. Weld (UW), Dr. Wei-Peng Chen (Fujitsu Lab), Prof. Huan Sun (OSU).



Mapping between natural language and programming language

Question

“how to clone or copy a python list?”



Code

```
“new_list =  
copy.copy(old_list)”
```

e.g., automated code **search/annotation/generation**.



Challenges

- Lack of large-scale datasets for model development.
i.e., pairs of <natural language question, code snippet>

Challenges

- Lack of large-scale datasets for model development.
i.e., pairs of <natural language question, code snippet>
- And datasets are important:



This work: StaQC

StaQC: the largest dataset to date of **~148K** Python and **~120K** SQL “how-to-do-it”* Question-Code pairs!

Example:

Question

“how to clone or copy a python list?”



Code

```
“new_list =  
copy.copy(old_list)”
```

“how-to-do-it” questions [Souza et al., 2014; Defim et al., 2016]: *the questioner provides a scenario and asks how to implement it.*

This work: StaQC

StaQC: the largest dataset to date of **~148K** Python and **~120K** SQL “how-to-do-it”* Question-Code pairs!

Continuously growing in *size* and *diversity*  stackoverflow

Diversity of StaQC

- Containing multiple code solutions to the same question.
 - Question “How to limit a number to be within a specified range?”
 - 4 code solutions in StaQC:

```
def clamp(n, minn, maxn):  
    return max(min(maxn, n),  
minn)
```

(1)

```
clamp = lambda n, minn,  
maxn: max(min(maxn, n),  
minn)
```

(2)

```
def clamp(n, minn, maxn):  
    if n < minn:  
        return minn  
    elif n > maxn:  
        return maxn  
    else:  
        return n
```

(3)

```
n = minn if n < minn else maxn if n > maxn else n
```

(4)

Diversity of StaQC

- Containing different questions asking for semantically similar code solutions.

Question A: *"How to find a gap in range in SQL"*

```
SELECT id + 1
FROM test mo
WHERE NOT EXISTS
(
    SELECT NULL
    FROM test mi
    WHERE mi.id = mo.id + 1
) and mo.id > 100
ORDER BY
    id
LIMIT 1
```

Question B: *"How do I find a "gap" in running counter with SQL?"*

```
SELECT id + 1
FROM mytable mo
WHERE NOT EXISTS
(
    SELECT NULL
    FROM mytable mi
    WHERE mi.id = mo.id + 1
)
ORDER BY
    id
LIMIT 1
```


Diversity of StaQC

- Containing different questions asking for semantically similar

Critical for Model Robustness:

1. Natural language variation.
2. Different implementations to do the same thing in programming language.

```
) and mo.id> 100  
ORDER BY  
      id  
LIMIT 1
```

```
)  
ORDER BY  
      id  
LIMIT 1
```

This work: StaQC

StaQC: the largest dataset to date of **~148K** Python and **~120K** SQL “how-to-do-it”* Question-Code pairs!

Continuously growing in size and diversity  stackoverflow

A better source for constructing models mapping
between NL and PL

This work: StaQC

StaQC: the largest dataset to date of **~148K** Python and **~120K** SQL “how-to-do-it”* Question-Code pairs!

Continuously growing in size and *diversity*



A better source for constructing models mapping
between NL and PL



Example



Question Elegant Python function to convert CamelCase to snake_case?

Accepted answer post



This is pretty thorough:

452



```
def convert(name):  
    s1 = re.sub('([A-Z][a-z]+)', r'\1_\2', name)  
    return re.sub('([a-z0-9])([A-Z])', r'\1_\2', s1).lower()
```

Works with all these (and doesn't harm already-un-cameled versions):

```
>>> convert('CamelCase')  
'camel_case'  
>>> convert('CamelCamelCase')  
'camel_camel_case'  
>>> convert('Camel2Camel2Case')  
'camel2_camel2_case'  
>>> convert('getHTTPResponseCode')  
'get_http_response_code'  
>>> convert('get2HTTPResponseCode')  
'get2_http_response_code'  
>>> convert('HTTPResponseCode')  
'http_response_code'  
>>> convert('HTTPResponseCodeXYZ')  
'http_response_code_xyz'
```

Or if you're going to call it a zillion times, you can pre-compile the regexes:

```
first_cap_re = re.compile('([A-Z][a-z]+)')  
all_cap_re = re.compile('([a-z0-9])([A-Z])')  
def convert(name):  
    s1 = first_cap_re.sub(r'\1_\2', name)  
    return all_cap_re.sub(r'\1_\2', s1).lower()
```

Example



Question Elegant Python function to convert CamelCase to snake_case?

Accepted answer post



This is pretty thorough:

452



```
def convert(name):  
    s1 = re.sub('([A-Z][a-z]+)', r'\1_\2', name)  
    return re.sub('([a-z0-9])([A-Z])', r'\1_\2', s1).lower()
```

Code block 1



Works with all these (and doesn't harm already-un-cameled versions):

```
>>> convert('CamelCase')  
'camel_case'  
>>> convert('CamelCamelCase')  
'camel_camel_case'  
>>> convert('Camel2Camel2Case')  
'camel2_camel2_case'  
>>> convert('getHTTPResponseCode')  
'get_http_response_code'  
>>> convert('get2HTTPResponseCode')  
'get2_http_response_code'  
>>> convert('HTTPResponseCode')  
'http_response_code'  
>>> convert('HTTPResponseCodeXYZ')  
'http_response_code_xyz'
```

Code block 2

Or if you're going to call it a zillion times, you can pre-compile the regexes:

```
first_cap_re = re.compile('([A-Z][a-z]+)')  
all_cap_re = re.compile('([a-z0-9])([A-Z])')  
def convert(name):  
    s1 = first_cap_re.sub(r'\1_\2', name)  
    return all_cap_re.sub(r'\1_\2', s1).lower()
```

Code block 3

Example



Question Elegant Python function to convert CamelCase to snake_case?

Accepted answer post

▲ This is pretty thorough:

452

```
def convert(name):
    s1 = re.sub('([A-Z][a-z]+)', r'\1_\2', name)
    return re.sub('([a-z0-9])([A-Z])', r'\1_\2', s1).lower()
```

Code block 1

Should we pair <Question, Code block n>?

```
>>> convert('Camel2Camel2Case')
'camel2_camel2_case'
>>> convert('getHTTPResponseCode')
'get_http_response_code'
>>> convert('get2HTTPResponseCode')
'get2_http_response_code'
>>> convert('HTTPResponseCode')
'http_response_code'
>>> convert('HTTPResponseCodeXYZ')
'http_response_code_xyz'
```

Code block 2

Or if you're going to call it a zillion times, you can pre-compile the regexes:

```
first_cap_re = re.compile('([A-Z][a-z]+)')
all_cap_re = re.compile('([a-z0-9])([A-Z])')
def convert(name):
    s1 = first_cap_re.sub(r'\1_\2', name)
    return all_cap_re.sub(r'\1_\2', s1).lower()
```

Code block 3

Example



Question Elegant Python function to convert CamelCase to snake_case?

Accepted answer post

▲ This is pretty thorough:

452

```
def convert(name):
    s1 = re.sub('([A-Z][a-z]+)', r'\1_\2', name)
    return re.sub('([a-z0-9])([A-Z])', r'\1_\2', s1).lower()
```

Code block 1

Is Code block n a “standalone” solution to the question?

```
>>> convert('Camel2Camel2Case')
'camel2_camel2_case'
>>> convert('getHTTPResponseCode')
'get_http_response_code'
>>> convert('get2HTTPResponseCode')
'get2_http_response_code'
>>> convert('HTTPResponseCode')
'http_response_code'
>>> convert('HTTPResponseCodeXYZ')
'http_response_code_xyz'
```

Code block 2

Or if you're going to call it a zillion times, you can pre-compile the regexes:

```
first_cap_re = re.compile('([A-Z][a-z]+)')
all_cap_re = re.compile('([a-z0-9])([A-Z])')
def convert(name):
    s1 = first_cap_re.sub(r'\1_\2', name)
    return all_cap_re.sub(r'\1_\2', s1).lower()
```

Code block 3

“Standalone” code solution

By looking at **Code block n**,

```
>>> convert('CamelCase')
'camel_case'
>>> convert('CamelCamelCase')
'camel_camel_case'
>>> convert('Camel2Camel2Case')
'camel2_camel2_case'
>>> convert('getHTTPResponseCode')
'get_http_response_code'
>>> convert('get2HTTPResponseCode')
'get2_http_response_code'
>>> convert('HTTPResponseCode')
'http_response_code'
>>> convert('HTTPResponseCodeXYZ')
'http_response_code_xyz'
```

can you solve the problem:

Elegant Python function to convert CamelCase to snake_case?

“Standalone” code solution

By looking at **Code block n**,

```
>>> convert('CamelCase')
'camel_case'
>>> convert('CamelCamelCase')
'camel_camel_case'
>>> convert('Camel2Camel2Case')
'camel2_camel2_case'
>>> convert('getHTTPResponseCode')
'get_http_response_code'
>>> convert('get2HTTPResponseCode')
'get2_http_response_code'
>>> convert('HTTPResponseCode')
'http_response_code'
>>> convert('HTTPResponseCodeXYZ')
'http_response_code_xyz'
```

can you solve the problem:

Elegant Python function to convert CamelCase to snake_case?

No! (it shows the usage, but with no details of the function)

Example: Question-Code pairs

Question Elegant Python function to convert CamelCase to snake_case?

Accepted answer post

This is pretty thorough:



```
def convert(name):
    s1 = re.sub('([A-Z][a-z]+)', r'\1_\2', name)
    return re.sub('([a-z0-9])([A-Z])', r'\1_\2', s1).lower()
```

**Standalone
solution!**

Works with all these (and doesn't harm already-un-cameled versions):



```
>>> convert('CamelCase')
'camel_case'
>>> convert('CamelCamelCase')
'camel_camel_case'
>>> convert('Camel2Camel2Case')
'camel2_camel2_case'
>>> convert('getHTTPResponseCode')
'get_http_response_code'
>>> convert('get2HTTPResponseCode')
'get2_http_response_code'
>>> convert('HTTPResponseCode')
'http_response_code'
>>> convert('HTTPResponseCodeXYZ')
'http_response_code_xyz'
```

**Not a *standalone*
solution!**
**(Showing usage, no
details of the function)**

Or if you're going to call it a zillion times, you can pre-compile the regexes:



```
first_cap_re = re.compile('([A-Z][a-z]+)')
all_cap_re = re.compile('([a-z0-9])([A-Z])')
def convert(name):
    s1 = first_cap_re.sub(r'\1_\2', name)
    return all_cap_re.sub(r'\1_\2', s1).lower()
```

**Standalone
solution!**

Previous methods: heuristics based

- **“Select All”**: Taking all code snippets in the answer post as code solutions. [Allamanis et al., 2015][Zilberstein and Yahav, 2016]
 - Low precision

(ground truth)

This is pretty thorough:



```
def convert(name):  
    s1 = re.sub('([A-Z][a-z]+)', r'\1_\2', name)  
    return re.sub('([a-z0-9])([A-Z])', r'\1_\2', s1).lower()
```

Works with all these (and doesn't harm already-un-cameled versions):



```
>>> convert('CamelCase')  
'camel_case'  
>>> convert('CamelCamelCase')  
'camel_camel_case'  
>>> convert('Camel2Camel2Case')  
'camel2_camel2_case'  
>>> convert('getHTTPResponseCode')  
'get_http_response_code'  
>>> convert('get2HTTPResponseCode')  
'get2_http_response_code'  
>>> convert('HTTPResponseCode')  
'http_response_code'  
>>> convert('HTTPResponseCodeXYZ')  
'http_response_code_xyz'
```

Or if you're going to call it a zillion times, you can pre-compile the regexes:



```
first_cap_re = re.compile('([A-Z][a-z]+)')  
all_cap_re = re.compile('([a-z0-9])([A-Z])')  
def convert(name):  
    s1 = first_cap_re.sub(r'\1_\2', name)  
    return all_cap_re.sub(r'\1_\2', s1).lower()
```

Previous methods: heuristics based

- **“Select First”**: Taking only the first code snippet in the answer post as a code solution, or considering only answer posts containing exactly one code snippet.

[Iyer et al., 2016]

□ **Low recall**

(ground truth)

This is pretty thorough:



```
def convert(name):  
    s1 = re.sub('([A-Z][a-z]+)', r'\1_\2', name)  
    return re.sub('([a-z0-9])([A-Z])', r'\1_\2', s1).lower()
```

Works with all these (and doesn't harm already-un-cameled versions):



```
>>> convert('CamelCase')  
'camel_case'  
>>> convert('CamelCamelCase')  
'camel_camel_case'  
>>> convert('Camel2Camel2Case')  
'camel2_camel2_case'  
>>> convert('getHTTPResponseCode')  
'get_http_response_code'  
>>> convert('get2HTTPResponseCode')  
'get2_http_response_code'  
>>> convert('HTTPResponseCode')  
'http_response_code'  
>>> convert('HTTPResponseCodeXYZ')  
'http_response_code_xyz'
```

Or if you're going to call it a zillion times, you can pre-compile the regexes:



```
first_cap_re = re.compile('([A-Z][a-z]+)')  
all_cap_re = re.compile('([a-z0-9])([A-Z])')  
def convert(name):  
    s1 = first_cap_re.sub(r'\1_\2', name)  
    return all_cap_re.sub(r'\1_\2', s1).lower()
```

Our solution: A systematic framework

■ Binary classification formulation:

Input: A question on Stack Overflow and its accepted answer post with multiple code snippets

Output: A binary label for each code snippet on whether it is a **standalone solution** to the question

This is pretty thorough:



```
def convert(name):  
    s1 = re.sub('([A-Z][a-z]+)', r'\1_\2', name)  
    return re.sub('([a-z0-9])([A-Z])', r'\1_\2', s1).lower()
```

Works with all these (and doesn't harm already-un-cameled versions):



```
>>> convert('CamelCase')  
'camel_case'  
>>> convert('CamelCamelCase')  
'camel_camel_case'  
>>> convert('Camel2Camel2Case')  
'camel2_camel2_case'  
>>> convert('getHTTPResponseCode')  
'get_http_response_code'  
>>> convert('get2HTTPResponseCode')  
'get2_http_response_code'  
>>> convert('HTTPResponseCode')  
'http_response_code'  
>>> convert('HTTPResponseCodeXYZ')  
'http_response_code_xyz'
```

Or if you're going to call it a zillion times, you can pre-compile the regexes:



```
first_cap_re = re.compile('([A-Z][a-z]+)')  
all_cap_re = re.compile('([a-z0-9])([A-Z])')  
def convert(name):  
    s1 = first_cap_re.sub(r'\1_\2', name)  
    return all_cap_re.sub(r'\1_\2', s1).lower()
```

A bi-view formulation

Interleaving
text and code
blocks

S₁ This is pretty thorough:

C₁

```
def convert(name):  
    s1 = re.sub('([A-Z][a-z]+)', r'\1_\2', name)  
    return re.sub('([a-z0-9])([A-Z])', r'\1_\2', s1).lower()
```

S₂ Works with all these (and doesn't harm already-un-cameled versions):

C₂

```
>>> convert('CamelCase')  
'camel_case'  
>>> convert('CamelCamelCase')  
'camel_camel_case'  
>>> convert('Camel2Camel2Case')  
'camel2_camel2_case'  
>>> convert('getHTTPResponseCode')  
'get_http_response_code'  
>>> convert('get2HTTPResponseCode')  
'get2_http_response_code'  
>>> convert('HTTPResponseCode')  
'http_response_code'  
>>> convert('HTTPResponseCodeXYZ')  
'http_response_code_xyz'
```

S₃ Or if you're going to call it a zillion times, you can pre-compile the regexes:

C₃

```
first_cap_re = re.compile('([A-Z][a-z]+)')  
all_cap_re = re.compile('([a-z0-9])([A-Z])')  
def convert(name):  
    s1 = first_cap_re.sub(r'\1_\2', name)  
    return all_cap_re.sub(r'\1_\2', s1).lower()
```

Text-based view: contextual hints

S_1 This is pretty thorough:

C_1  more likely to be a code solution


S_2 Works with all these (and doesn't harm already-un-cameled versions):

C_2

S_3 Or if you're going to call it a zillion times you can pre-compile the regexes:

C_3  more likely to be a code solution

Code-based view: semantics of code content

C_1	<pre>def convert(name): s1 = re.sub('([A-Z][a-z]+)', r'\1_2', name) return re.sub('([a-z0-9])([A-Z])', r'\1_2', s1).lower()</pre>		more likely to be a solution
C_2	<pre>>>> convert('CamelCase') 'camel_case' >>> convert('CamelCamelCase') 'camel_camel_case' >>> convert('Camel2Camel2Case') 'camel2_camel2_case' >>> convert('getHTTPResponseCode') 'get_http_response_code' >>> convert('get2HTTPResponseCode') 'get2_http_response_code' >>> convert('HTTPResponseCode') 'http_response_code' >>> convert('HTTPResponseCodeXYZ') 'http_response_code_xyz'</pre>		possibly a usage demo
C_3	<pre>first_cap_re = re.compile('([A-Z][a-z]+)') all_cap_re = re.compile('([a-z0-9])([A-Z])') def convert(name): s1 = first_cap_re.sub(r'\1_2', name) return all_cap_re.sub(r'\1_2', s1).lower()</pre>		more likely to be a solution

Formulation for each code snippet

Predict a “solution or not” label for a code snippet (here C_2) based on:

1. **Textual context** (text view): S_2, S_3 .
2. **Code content** (code view): C_2 .

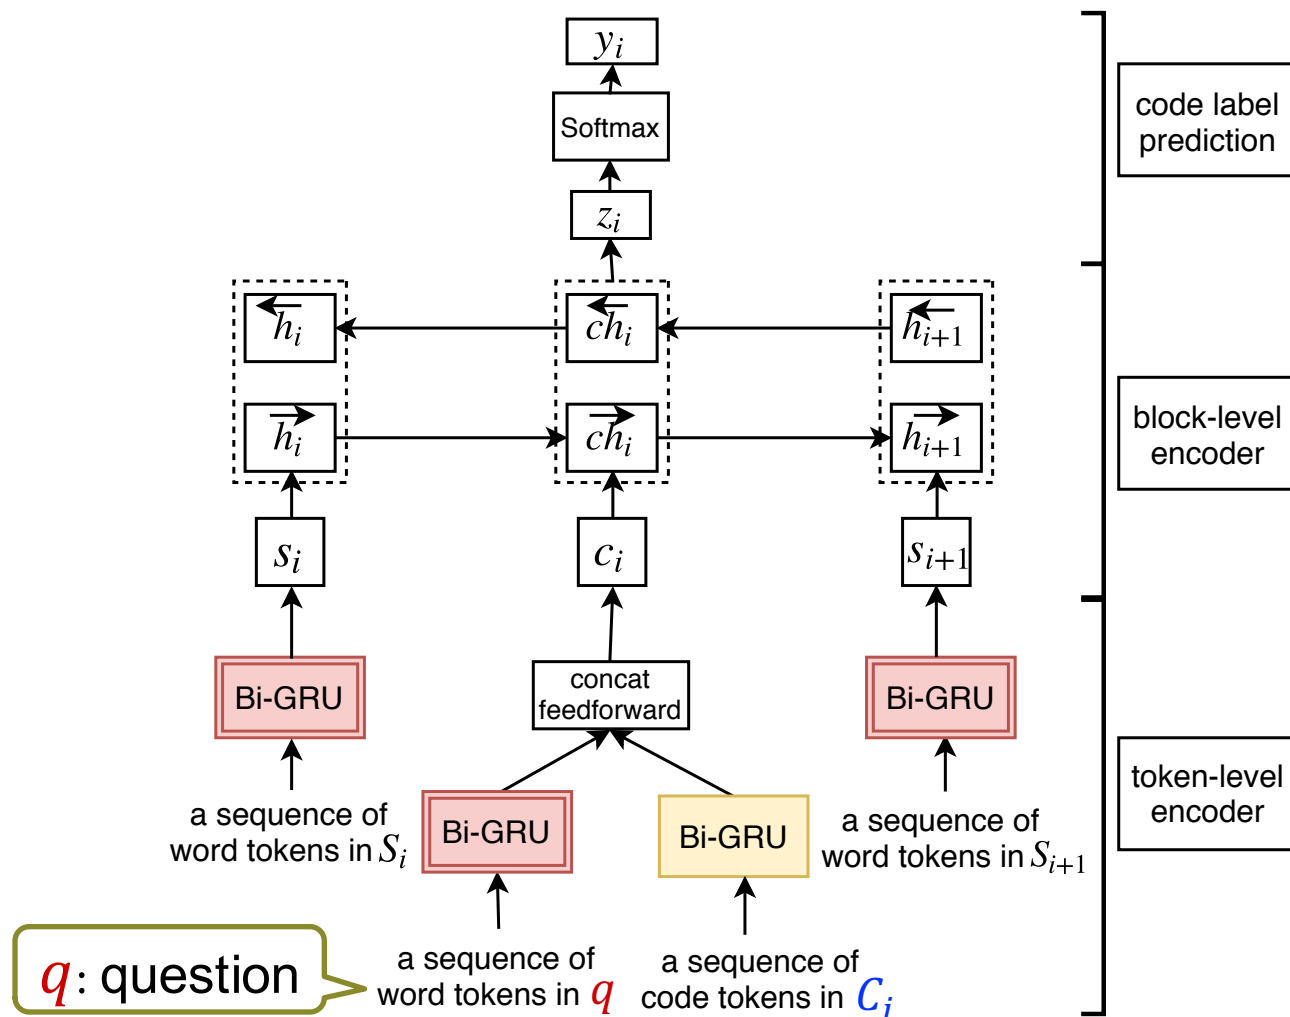
S_2 Works with all these (and doesn't harm already-un-cameled versions):

C_2

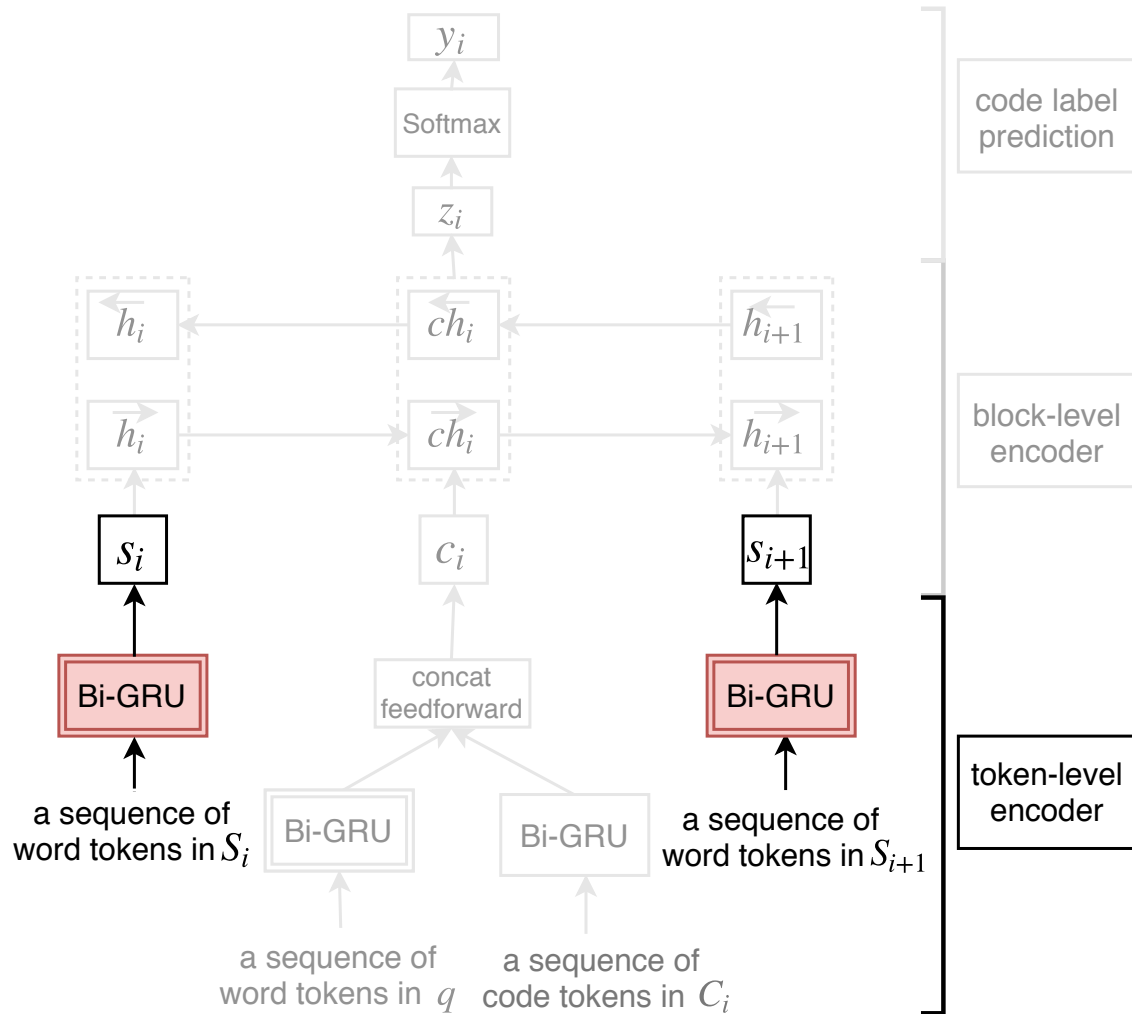
```
>>> convert('CamelCase')
'camel_case'
>>> convert('CamelCamelCase')
'camel_camel_case'
>>> convert('Camel2Camel2Case')
'camel2_camel2_case'
>>> convert('getHTTPResponseCode')
'get_http_response_code'
>>> convert('get2HTTPResponseCode')
'get2_http_response_code'
>>> convert('HTTPResponseCode')
'http_response_code'
>>> convert('HTTPResponseCodeXYZ')
'http_response_code_xyz'
```

S_3 Or if you're going to call it a zillion times, you can pre-compile the regexes:

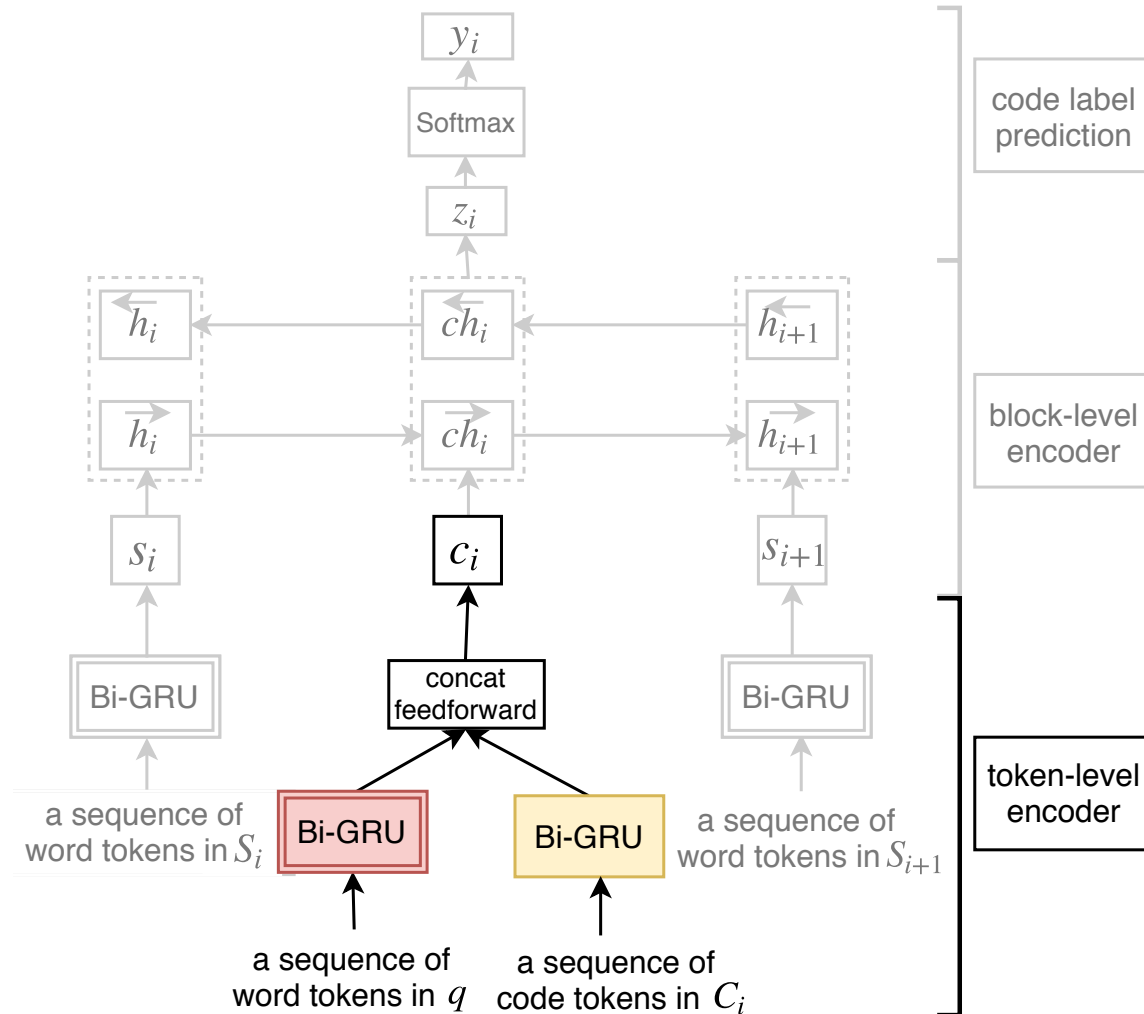
Bi-View Hierarchical Neural Network (BiV-HNN)



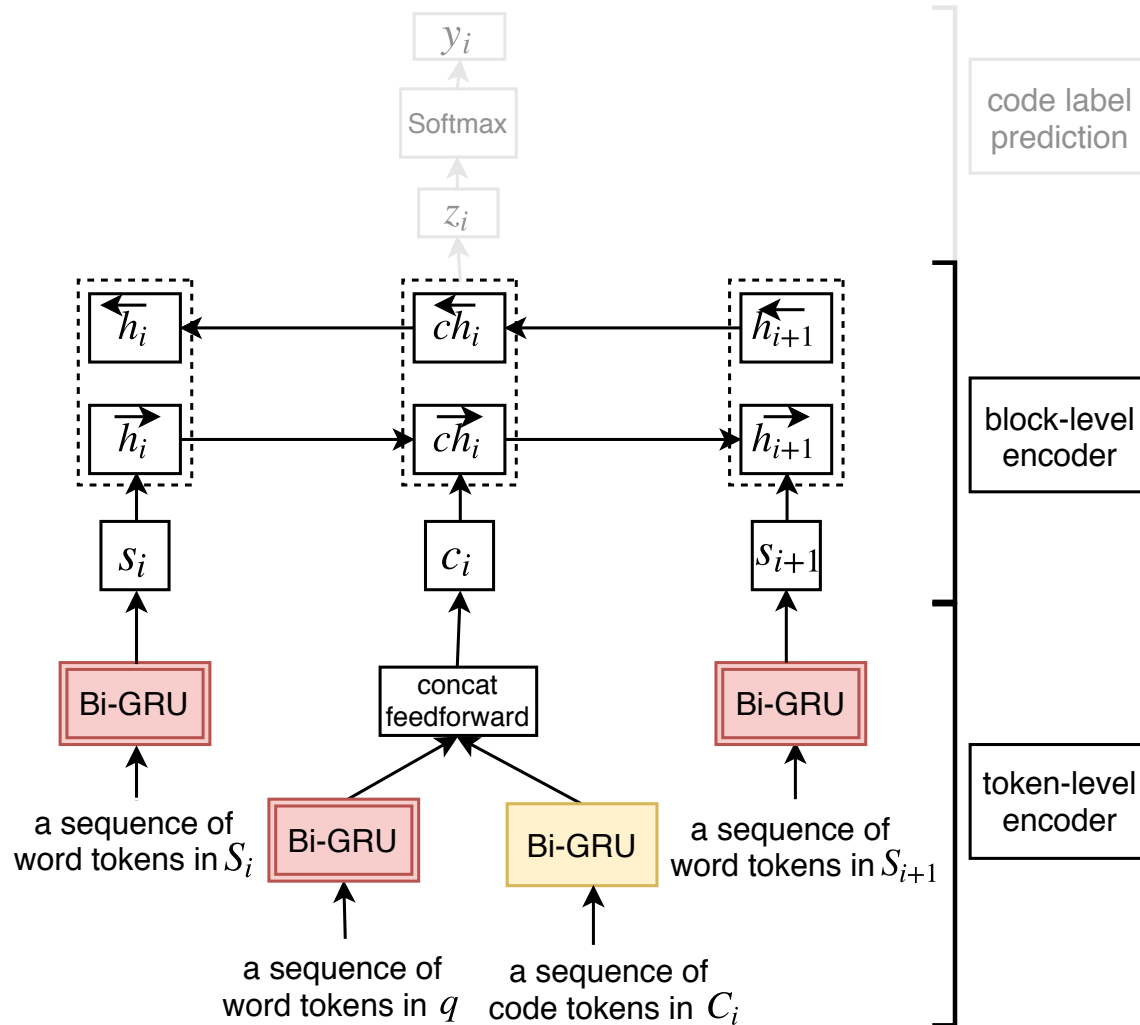
Token-level encoder for text blocks



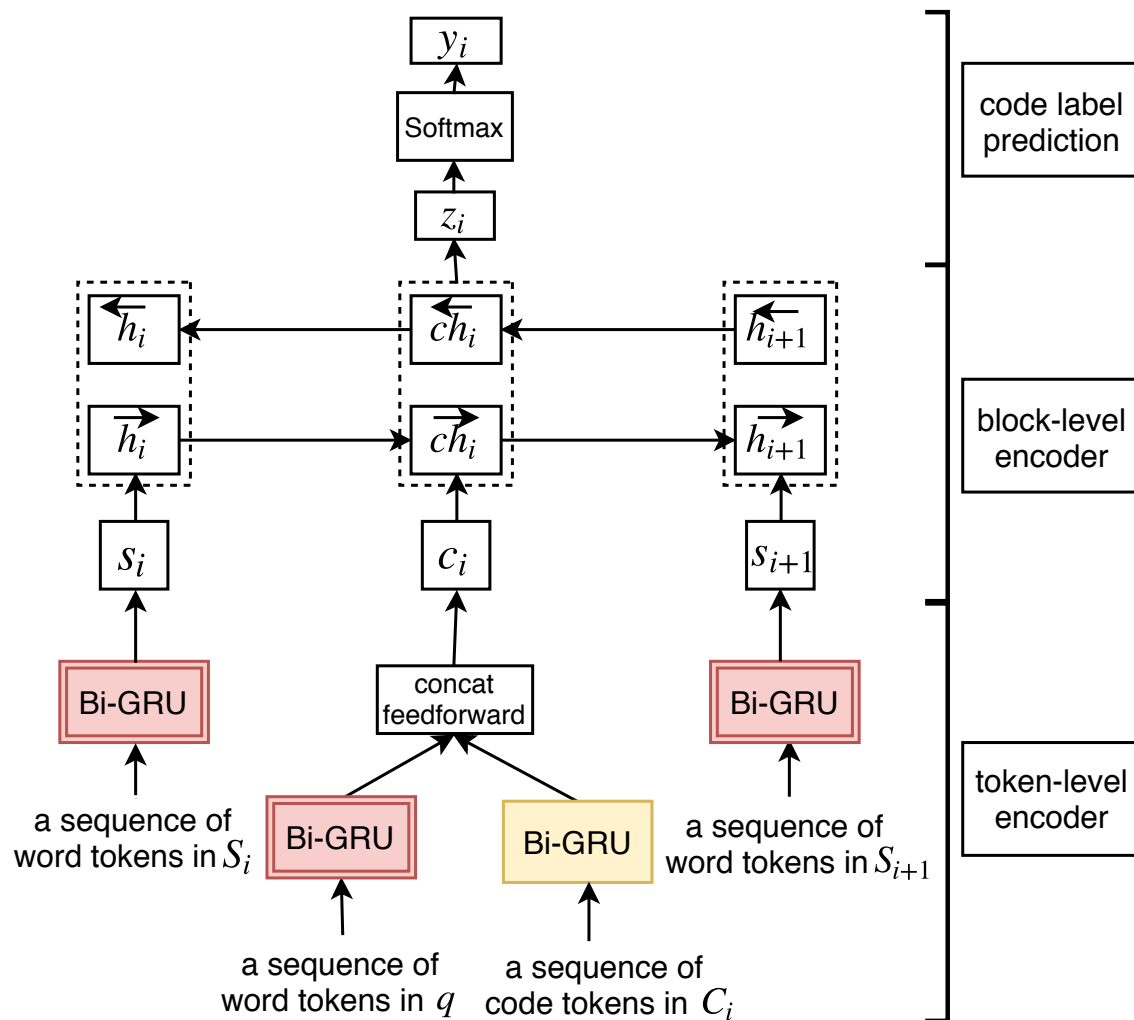
Token-level encoder for code blocks



Block-level encoder



Code label prediction



Experimental Setup

- Manually annotating “solution or not” label on code snippets in one answer post.
 - Python and SQL domain.
 - Four undergraduates with substantial Cohen’s kappa agreement.
- Training/validation/test split: 60% - 20% -20%.

	Python	SQL
# of Question-Code pairs	4,884	3,637
% of positive Question-Code pairs	44%	57%

Main Results

- Heuristic methods: Select-First, Select-All.
- Feature engineering based methods:
 - Logistic Regression (LR), Support Vector Machine (SVM).
 - Features: text-based (uni-/bi-grams, the connectives, etc) and code-based (code tokens, etc).

	Python	SQL
Select-First	0.607	0.613
Select-All	0.642	0.737
LR	0.766	0.846
SVM	0.753	0.850
BiV-HNN	0.841	0.888

(comparison on F1)

Research questions for understanding BiV-HNN

Q1: text view, code view, or bi-view?

Q2: hierarchical structure or flat structure?

Q3: block-level encoder: sequential or feedforward?


Research questions for understanding BiV-HNN

Q1: text view, code view, or bi-view?

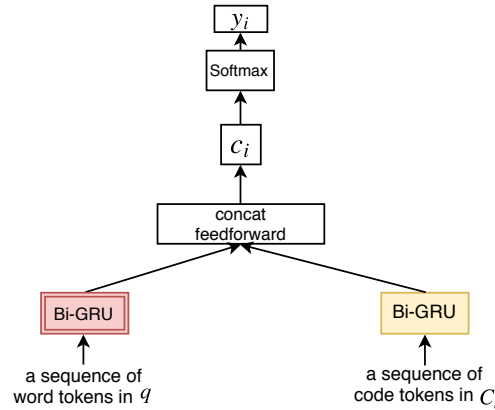
Q2: hierarchical structure or flat structure?

Q3: block-level encoder: sequential or feedforward?

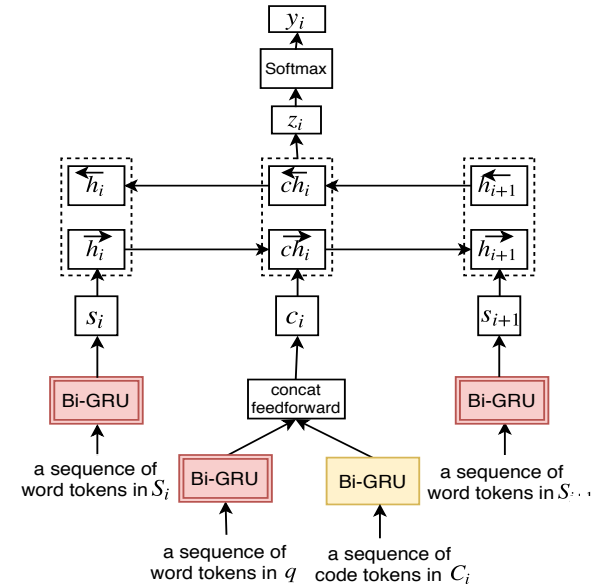
Q1: text view, code view, or bi-view?



Text-HNN
(text view)



Code-HNN
(code view)



BiV-HNN
(bi-view)

	Python	SQL
Text-HNN	0.771	0.840
Code-HNN	0.812	0.851
BiV-HNN	0.841	0.888

(comparison on F1)

Model combination

- Text-HNN, Code-HNN and BiV-HNN are observed **complementary** to each other.
 - On Python validation set, 60%~70% of mistakes made by one model can be corrected by the other two models.

Model combination

- Text-HNN, Code-HNN and BiV-HNN are observed **complementary** to each other.
 - On Python validation set, 60%~70% of mistakes made by one model can be corrected by the other two models.
- **Model combination**: the label of a code snippet is predicted only when the three models agree on it.

Model combination

- Text-HNN, Code-HNN and BiV-HNN are observed **complementary** to each other.
 - On Python validation set, 60%~70% of mistakes made by one model can be corrected by the other two models.
- **Model combination**: the label of a code snippet is predicted only when the three models agree on it.
- Model combination on testing set:
 - Python: ~70% of code snippets are labeled with **0.92 F1**.
 - SQL: ~80% of code snippets are labeled with **0.94 F1**.

Systematically mined StaQC

This is pretty thorough:

```
def convert(name):
    s1 = re.sub('([A-Z][a-z]+)', r'\1_2', name)
    return re.sub('([a-z0-9])([A-Z])', r'\1_2', s1).lower()
```


Works with all these (and doesn't harm already-un-cameled versions):

```
>>> convert('CamelCase')
'camel_case'
>>> convert('CamelCamelCase')
'camel_camel_case'
>>> convert('Camel2Camel2Case')
'camel2_camel2_case'
>>> convert('getHTTPResponseCode')
'get_http_response_code'
>>> convert('get2HTTPResponseCode')
'get2_http_response_code'
>>> convert('HTTPResponseCode')
'http_response_code'
>>> convert('HTTPResponseCodeXYZ')
'http_response_code_xyz'
```

Or if you're going to call it a zillion times, you can pre-compile the regexes:

```
first_cap_re = re.compile('([A-Z][a-z]+)')
all_cap_re = re.compile('([a-z0-9])([A-Z])')
def convert(name):
    s1 = first_cap_re.sub(r'\1_2', name)
    return all_cap_re.sub(r'\1_2', s1).lower()
```

New or unannotated post



**Model
combination**

This is pretty thorough:



```
def convert(name):
    s1 = re.sub('([A-Z][a-z]+)', r'\1_2', name)
    return re.sub('([a-z0-9])([A-Z])', r'\1_2', s1).lower()
```

Works with all these (and doesn't harm already-un-cameled versions):



```
>>> convert('CamelCase')
'camel_case'
>>> convert('CamelCamelCase')
'camel_camel_case'
>>> convert('Camel2Camel2Case')
'camel2_camel2_case'
>>> convert('getHTTPResponseCode')
'get_http_response_code'
>>> convert('get2HTTPResponseCode')
'get2_http_response_code'
>>> convert('HTTPResponseCode')
'http_response_code'
>>> convert('HTTPResponseCodeXYZ')
'http_response_code_xyz'
```

Or if you're going to call it a zillion times, you can pre-compile the regexes:



```
first_cap_re = re.compile('([A-Z][a-z]+)')
all_cap_re = re.compile('([a-z0-9])([A-Z])')
def convert(name):
    s1 = first_cap_re.sub(r'\1_2', name)
    return all_cap_re.sub(r'\1_2', s1).lower()
```

StaQC: Systematically mined Question-Code pairs

StaQC: the largest dataset to date of **~148K** Python and **~120K** SQL “how-to-do-it”* Question-Code pairs!

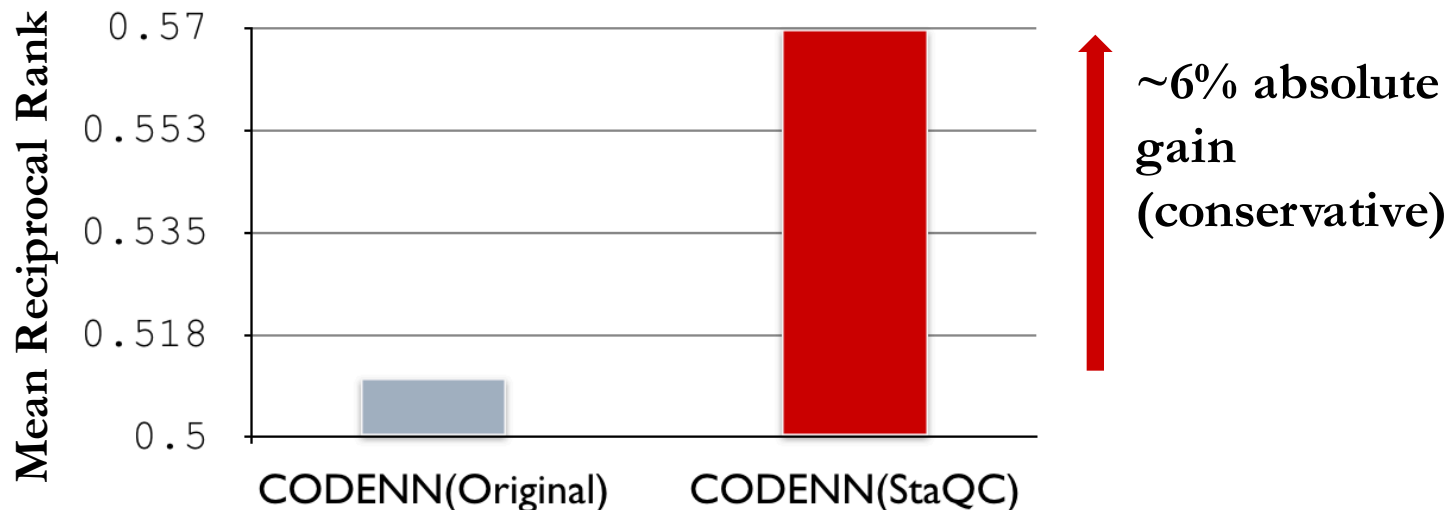
Continuously growing in size and diversity



A better source for constructing models mapping between natural language and programming language

StaQC: A better source for downstream tasks

- **Code Retrieval** as an exemplar downstream task (SQL domain).
- **Neural Network model CODENN.** [Iyer et al., 2016]
 - ▣ CODENN(Original): trained on ~26K heuristically collected QC pairs.
 - ▣ CODENN(StaQC): trained on ~120K systematically mined QC pairs.



StaQC: Systematically mined Question-Code pairs

StaQC: the largest dataset to date of **~148K** Python and **~120K** SQL “how-to-do-it”* Question-Code pairs!

Continuously growing in size and *diversity*  stackoverflow

A better source for constructing models mapping between natural language and programming language

Data and code are available at:

<https://github.com/LittleYUYU/StackOverflow-Question-Code-Dataset>

Thank you! Questions?